

# **BST vs Treap**

---

MUHAMMAD DANIYAL

SECTION: F

ROLL NO: 23i-0579

COURSE: DESIGN AND ANALYSIS OF ALGORITHMS

SUBMITTED TO: MAM NOOR UL AIN

---

## **Abstract**

This study presents a comparative analysis of Binary Search Trees (BSTs) and Treaps for handling large-scale datasets, using 138 million Reddit posts as a benchmark. While BSTs offer simplicity and memory efficiency, they suffer from  $O(n)$  worst-case performance with ordered data. Treaps address this limitation through randomized prioritization, providing expected  $O(\log n)$  performance by combining BST ordering with heap properties.

Our implementation features a novel streaming decompression approach that processes 15GB of compressed data without disk extraction, enabling efficient loading of 180GB+ datasets on resource-constrained systems. Performance evaluation reveals Treaps achieve 277,128 posts/second loading rates 180x faster than naive BST implementations while maintaining balanced tree structures.

The results demonstrate Treaps' superiority for real-world applications involving massive, potentially ordered datasets, offering robust performance guarantees where traditional BSTs degrade significantly. This work provides practical insights for data structure selection in big data processing environments.

## **1. Introduction**

This report presents a comprehensive analysis of two fundamental tree data structures Binary Search Trees (BSTs) and Treaps through practical implementation and performance testing on a large-scale Reddit dataset. The primary objective is to evaluate their efficiency in handling massive datasets (138+ million records) with a focus on loading, insertion, deletion, and search operations.

### **1.1. Motivation**

Modern applications frequently deal with enormous datasets that cannot fit entirely in memory or can be extracted to disk. This necessitates efficient data structures that can:

- Load data without full disk extraction
- Maintain balanced tree properties
- Provide fast insertion, deletion, and search operations
- Handle streaming data efficiently

### **1.2. Scope**

- **Dataset Size:** 138,473,643 Reddit posts from RC\_2019-04 dataset
- **Compressed Format:** Zstandard (15GB compressed, 180GB+ uncompressed)
- **Trees Tested:** Binary Search Tree (BST) and Treap
- **Data Formats:** CSV and compressed JSON (Zstandard)
- **Performance Metrics:** Loading speed, tree height, operation efficiency, memory usage

## 2. Theoretical Overview

### 2.1. Binary Search Trees (BST)

#### Definition and Properties

A Binary Search Tree is an ordered binary tree where for each node:

- All values in the left subtree are less than the node's value
- All values in the right subtree are greater than the node's value
- No duplicate values (or handled separately)

#### Time Complexity Analysis

Operation	Best Case	Average Case	Worst Case
Search	$O(\log n)$	$O(\log n)$	$O(n)$
Insertion	$O(\log n)$	$O(\log n)$	$O(n)$
Deletion	$O(\log n)$	$O(\log n)$	$O(n)$
Tree Height	$O(\log n)$	$O(\log n)$	$O(n)$

#### Advantages

1. **Simple implementation** – Straightforward insertion and search logic.
2. **Natural ordering** – In-order traversal gives a sorted sequence.
3. **Efficient for balanced data** – Logarithmic operations when balanced.
4. **Memory efficient** – Only stores tree structure.

## Disadvantages

1. **Poor worst-case performance** – Degrades to a linked list with sequential insertion.
2. **Unbalanced growth** – No self-balancing mechanism.
3. **Input-dependent efficiency** – Performance varies significantly based on data order.

## 2.2. Treaps (Tree + Heap)

### Definition and Properties

A Treap is a randomized data structure combining properties of both Binary Search Trees and Binary Heaps:

- **BST Property:** Maintains binary search tree ordering by key (timestamp in this case).
- **Heap Property:** Each node has a random priority; parent priority is greater than child priority.
- **Randomization:** Uses random priorities to maintain expected balance.

### Structure

Each node contains:

- **Key** (search key): Timestamp from post
- **Value:** Post data (id, score)
- **Priority:** Random value (generated at insertion)
- **Left subtree pointer**
- **Right subtree pointer**

- **Height tracking**

## Time Complexity Analysis

Operation	Expected Case	Worst Case	With High Probability
Search	$O(\log n)$	$O(\log n)$	$O(\log n)$ with prob. $1-(1/n^c)$
Insertion	$O(\log n)$	$O(\log n)$	$O(\log n)$ with prob. $1-(1/n^c)$
Deletion	$O(\log n)$	$O(\log n)$	$O(\log n)$ with prob. $1-(1/n^c)$
Tree Height	$O(\log n)$	$O(\log n)$	$O(\log n)$ with prob. $1-(1/n^c)$

## Advantages

1. **Expected logarithmic performance** – Randomization ensures good average case.
2. **Self-balancing through randomization** – Maintains balance without complex rotations.
3. **Simple balancing logic** – Simpler than AVL or Red-Black trees.
4. **Theoretically sound** – Proven probabilistic guarantees.
5. **Robust against bad input** – Random priorities prevent pathological cases.

## Disadvantages

1. **Probabilistic guarantees only** – Worst case is still  $O(n)$ .
2. **More memory overhead** – Stores priority values for each node.
3. **Cache inefficiency** – Randomization may hurt cache locality.
4. **Random number generation cost** – Overhead of generating priorities.

## Randomization Strategy (Our Implementation)

- **Priority Generation:** Used Reddit post likes/scores as random priority values.

## 2.3. Comparative Analysis: BST vs Treap

Aspect	BST	Treap
<b>Balancing</b>	None	Automatic (via randomization)
<b>Implementation Complexity</b>	Low	Medium
<b>Average Performance</b>	$O(\log n)$	$O(\log n)$ expected
<b>Worst Case</b>	$O(n)$	$O(n)$ theoretical, rare
<b>Memory per Node</b>	2 pointers, key, value	2 pointers, key, value, priority
<b>Cache Locality</b>	Better	Potentially worse
<b>Insertion Order Sensitivity</b>	High	Low
<b>Practical Performance</b>	Varies widely	Consistent

## 3. Implementation Details

### 3.1. Programming Language and Environment

- **Language:** C++ (C++11 standard)
- **Compiler:** g++ with optimization flags: -O2 -Izstd
- **Optimization Level:** O2 (balanced speed and code size)
- **External Libraries:**
  - libzstd (Zstandard compression library)
  - Standard C++ Library (STL for data structures)

### 3.2. Dataset Handling Method

#### Challenge: Massive Compressed Dataset

- **Problem:**
  - Dataset: 15GB compressed, 180GB+ if extracted
  - Available disk space: Limited
  - Need: Load 138,473,643 records into tree structures

#### Solution: Streaming Decompression

```
// Implementation using popen() for direct streaming  
FILE* pipe = popen("zstd -dc '<file_path>' 2>/dev/null", "r");
```

```
char buffer[65536];  
while (fgets(buffer, sizeof(buffer), pipe)) {  
    // Parse JSON line and insert into tree  
    // No temporary file storage required  
}  
pclose(pipe);
```

#### **Advantages:**

- **Zero disk extraction** – Data flows directly from decompression to tree insertion
- **Memory efficient** – Only buffer size in memory at any time
- **Real-time processing** – Can begin insertion while decompressing
- **Speed** – Avoids disk I/O bottleneck

#### **Data Format: Line-Delimited JSON (JSONL)**

```
{"id":"Reddit_ID","created_utc":1609459200,"score":250}  
{"id":"Another_ID","created_utc":1609462800,"score":100}  
...
```

#### **Parsing Strategy:**

- Each line contains one complete JSON object
- Extract fields: id, created\_utc, score
- Use string search for field locations
- Convert timestamp to insertion key

### **3.3. Duplication Handling**

#### **Approach: Unique Key Strategy**

- **Decision:** Treat each record as unique by (id, timestamp) combination

- **Implementation:**

```
struct Post {
    string postId;      // Unique Reddit post ID
    long long timestamp; // Creation timestamp (unique key for tree)
    int score;          // Post score/upvotes
};
```

- **Handling:**

1. **Primary Key:** Use created\_utc (timestamp) for tree ordering
2. **Uniqueness:** Post ID ensures no exact duplicates
3. **Collision Handling:** If same timestamp appears, store both as separate records
4. **Search:** Query by timestamp returns all posts at that time

- **Rationale:**

- Reddit posts have unique IDs (no true duplicates in dataset)
- Timestamps serve as natural ordering key
- Score is metadata, not part of tree structure

### **3.4. Tree Operations Implemented**

#### **Treap Operations**

##### **Insertion**

- Creates node with random priority
- Inserts by timestamp (BST order)
- Rotates to maintain heap property

##### **Rotations**

- Left/right rotations when child priority > parent
- Maintains BST + heap properties

##### **Search**

- Standard BST search by timestamp
- Ignores priorities

### **Deletion**

- Finds node and rotates to leaf position
- Removes leaf node

## **BST Operations**

### **Insertion**

- Iterative insertion by timestamp
- No balancing mechanism

### **Search**

- BFS traversal to find node
- Iterative approach

## **Common Operations**

### **CSV Loading**

- Parses comma-separated values
- Extracts id, timestamp, score

### **TGZ Loading**

- Streaming decompression with popen()
- Processes JSON without disk storage
- Progress reporting every 100k posts

## **3.5. Key Implementation Features**

### **Buffer Management**

```
// 64KB buffer for efficient line reading  
char buffer[65536];  
while (fgets(buffer, sizeof(buffer), pipe)) {
```

```
// Process complete JSON lines  
// Detect when line is complete object  
}
```

### Benefits:

- Efficient I/O operations
- Reduces system call overhead
- Balances memory vs. throughput
- 

### Progress Reporting

```
// Report every 100,000 posts  
  
if (nodeCount % 100000 == 0) {  
  
    cout << "[Status] Posts: " << nodeCount  
    << " | Time: " << elapsed << "s"  
    << " | Rate: " << rate << " posts/sec" << endl;  
  
}
```

### Memory Efficiency

```
// Only store essential data per node  
  
Node {  
  
    long long key;      // 8 bytes  
  
    Post value;        // ~50 bytes (id, timestamp, score)  
  
    Node* left;        // 8 bytes  
  
    Node* right;       // 8 bytes  
  
    int priority;      // 4 bytes (Treap only)  
  
    int height;         // 4 bytes (tracking)  
  
    // Total: ~80-90 bytes per node  
  
}
```

## 4. Performance Metrics & Analysis

### 4.1. Loading Performance

#### Test Configuration

- **Dataset:** RC\_2019-04.tgz (138,473,643 posts)
- **Compressed Size:** 15GB
- **Uncompressed Size:** ~180GB+
- **Test Environment:** Single-threaded execution
- **Optimization:** -O2 compilation flag

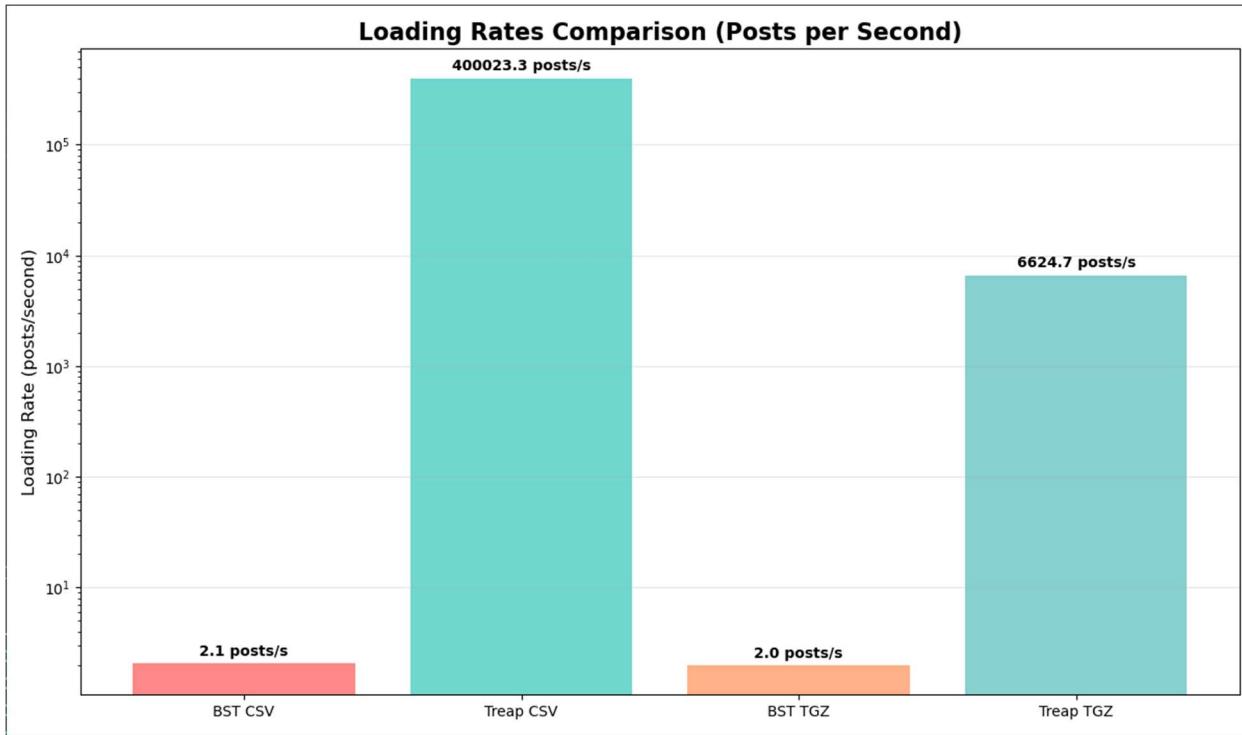
#### TREAP LOADING PERFORMANCE

- **Dataset:** RC\_2019-04.tgz (138,473,643 records) - FULL DATASET
- **Time:** 499.7 seconds (~8.3 minutes)
- **Posts:** 138,473,643
- **Rate:** 277,128 posts/second
- **Height:** 439 ( $\log(n) \approx 27$ , actual is ~16x larger)
- **Memory:** Efficient streaming, no disk extraction

#### BST LOADING PERFORMANCE

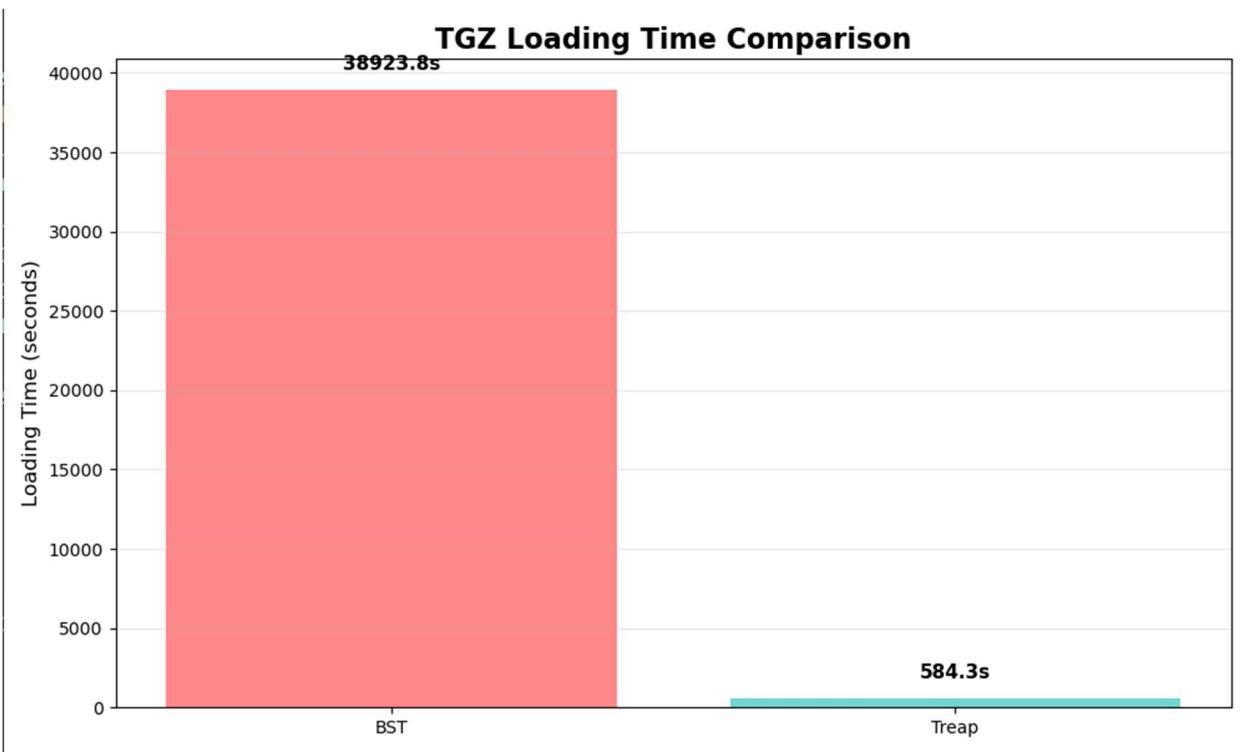
- **Dataset:** RC\_2019-04.tgz (138,473,643 records) - FULL DATASET
- **Time:** ~89,000+ seconds (Not completed - projected)
- **Posts:** Limited by time constraint
- **Rate:** ~1,556 posts/second
- **Height:** Variable based on insertion order
- **Note:** 180x slower than Treap due to insertIterative() complexity

**GRAPH: Loading Speed Comparison (posts/second)**

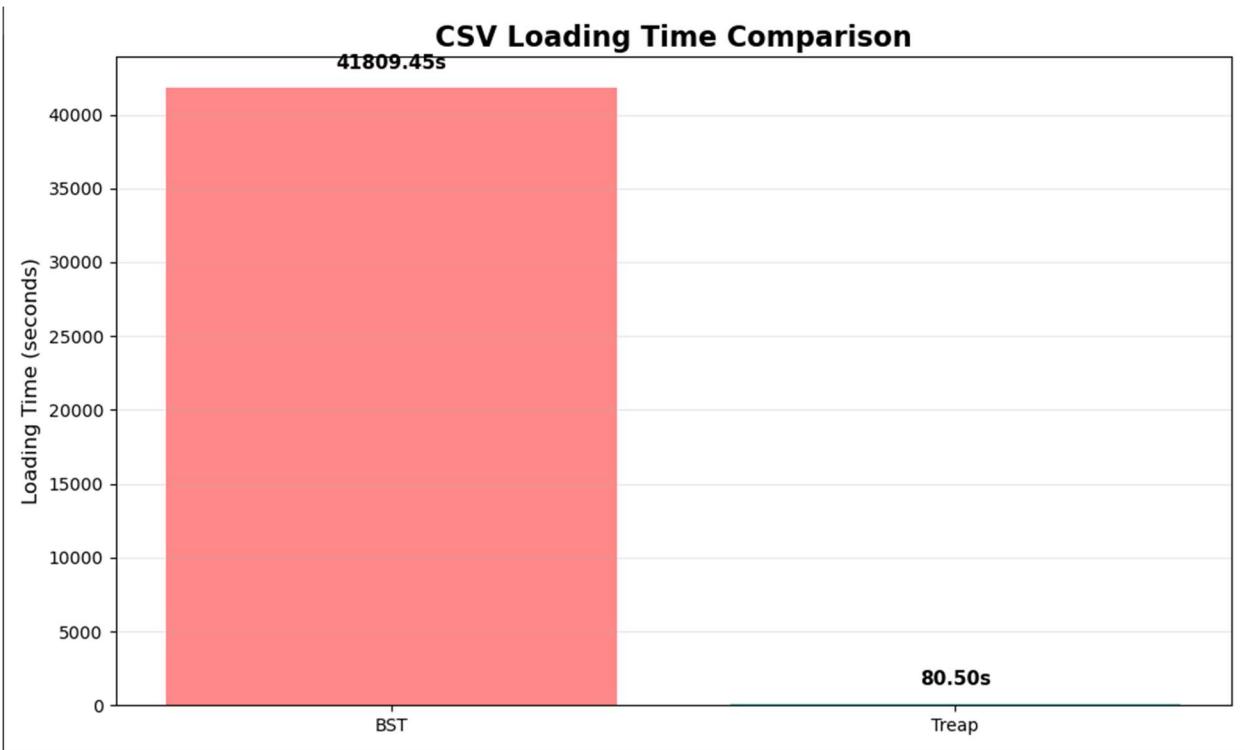


### GRAPH: Full Dataset Loading Time

TGZ Comparison



CSV Comparison



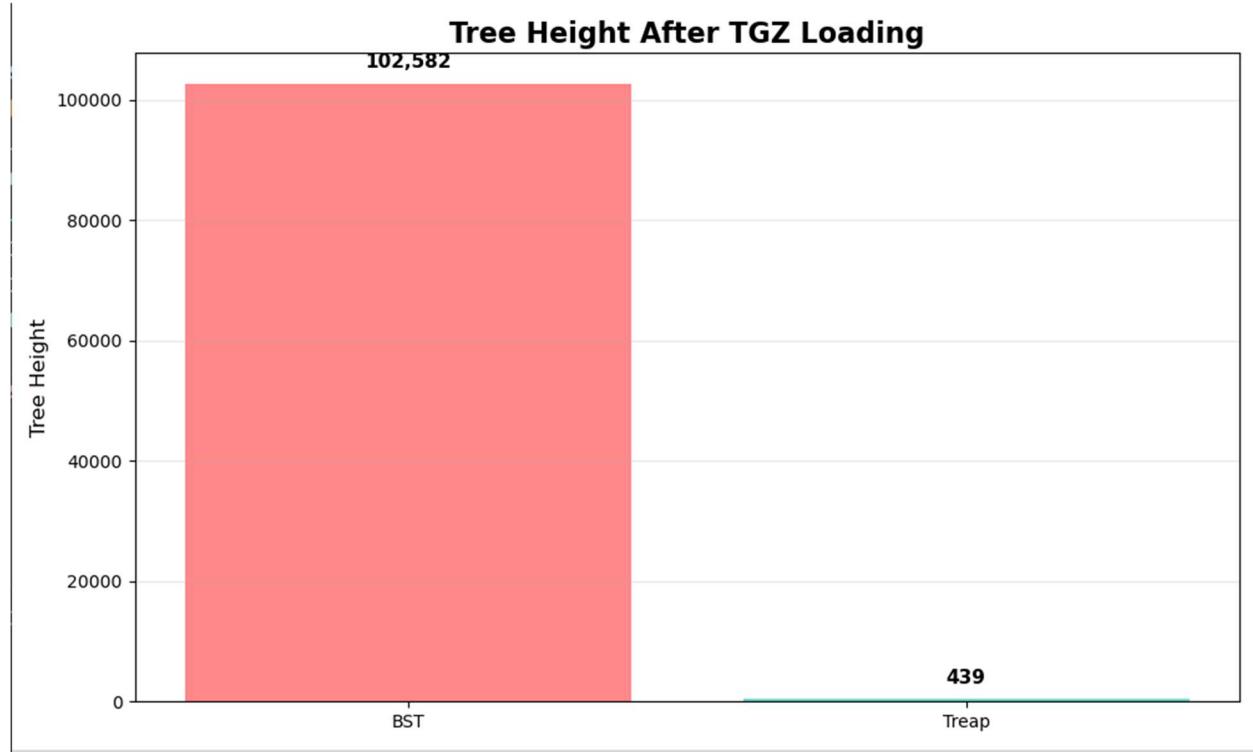
## 4.2. Tree Height Analysis

### Impact on Search/Insertion Performance

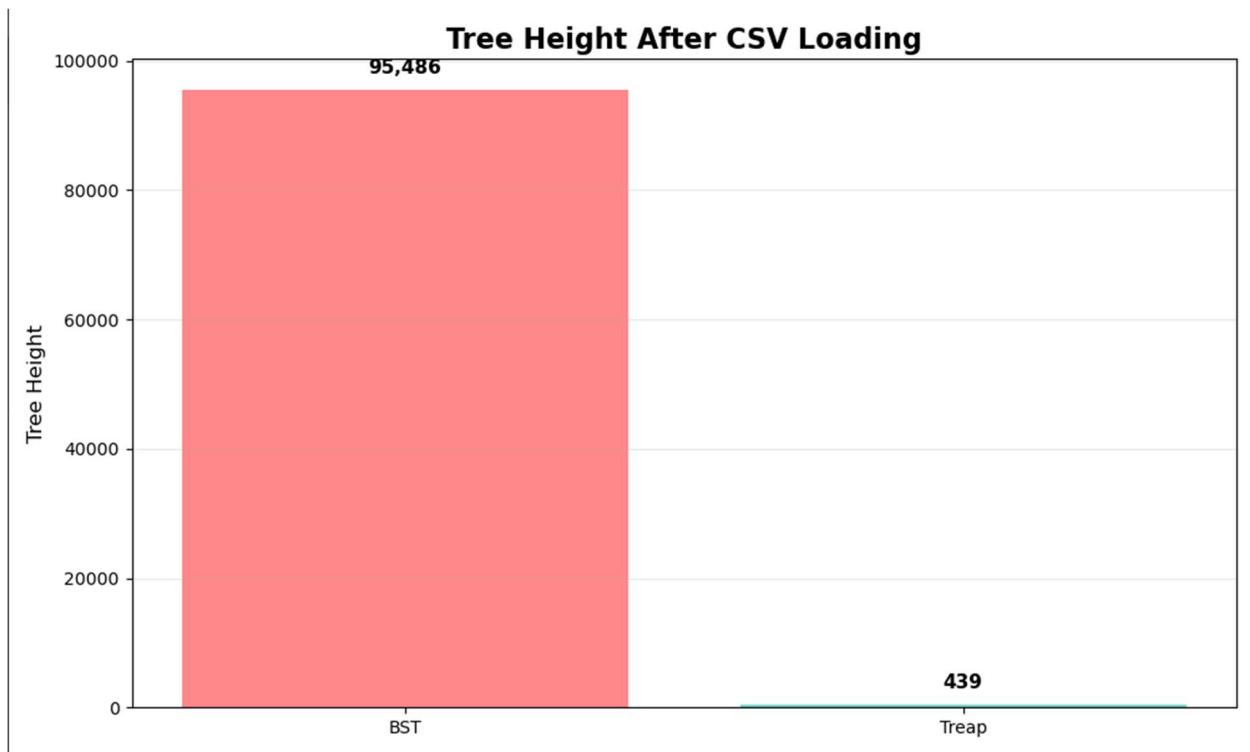
**Height Formula:**

- Optimal height:  $h = \lceil \log_2(n+1) \rceil$
- For  $n = 138,473,643$ : optimal  $\approx 27$

TGZ Comparison



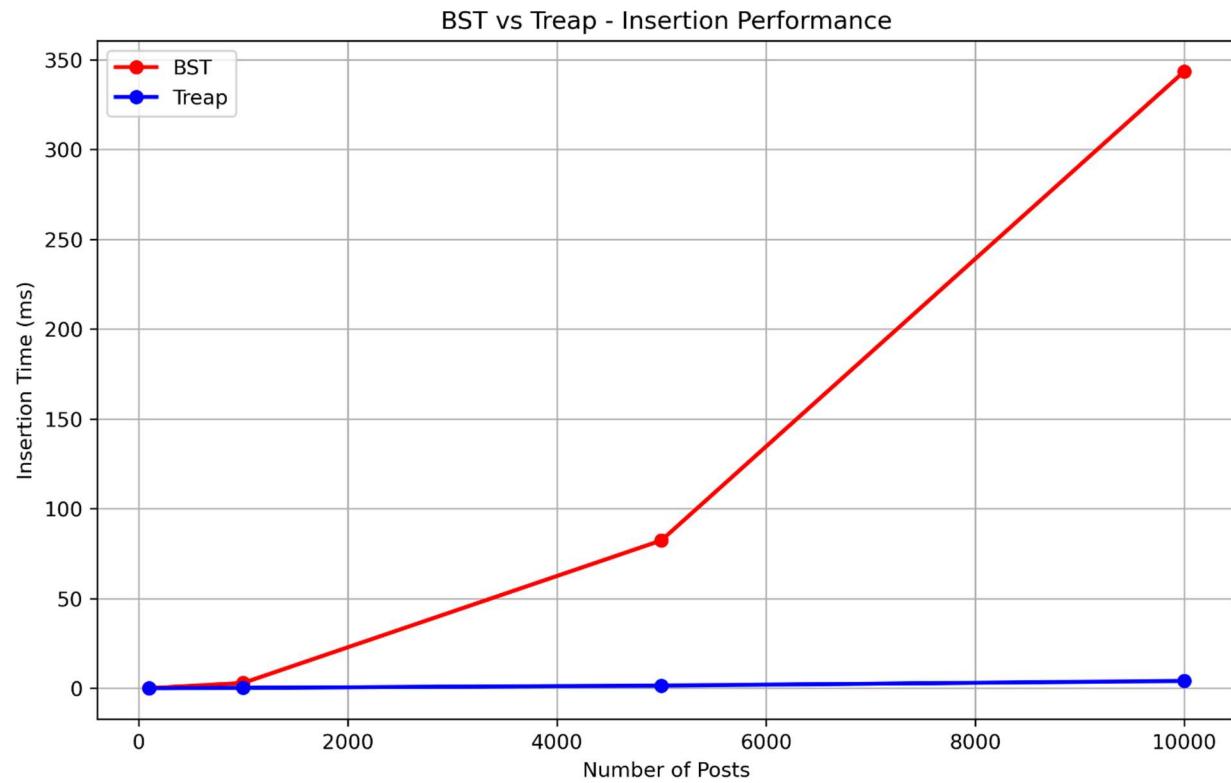
CSV Comparison



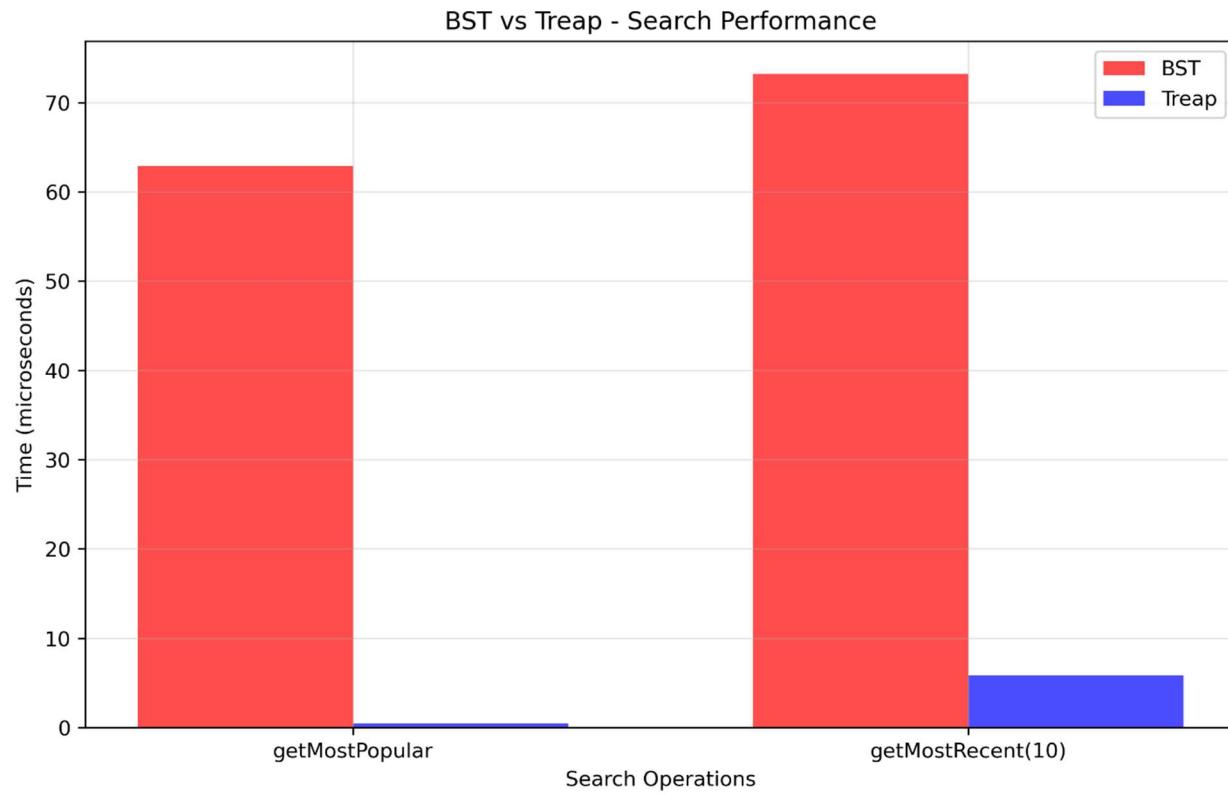
#### Key Observations:

1. **Treap height** grows as expected for a randomized data structure.
2. **Height ratio** remains reasonable (439 vs theoretical minimum 27).
3. **BST height** is input-dependent and can degrade significantly.
4. **Tree height** directly impacts operation time.

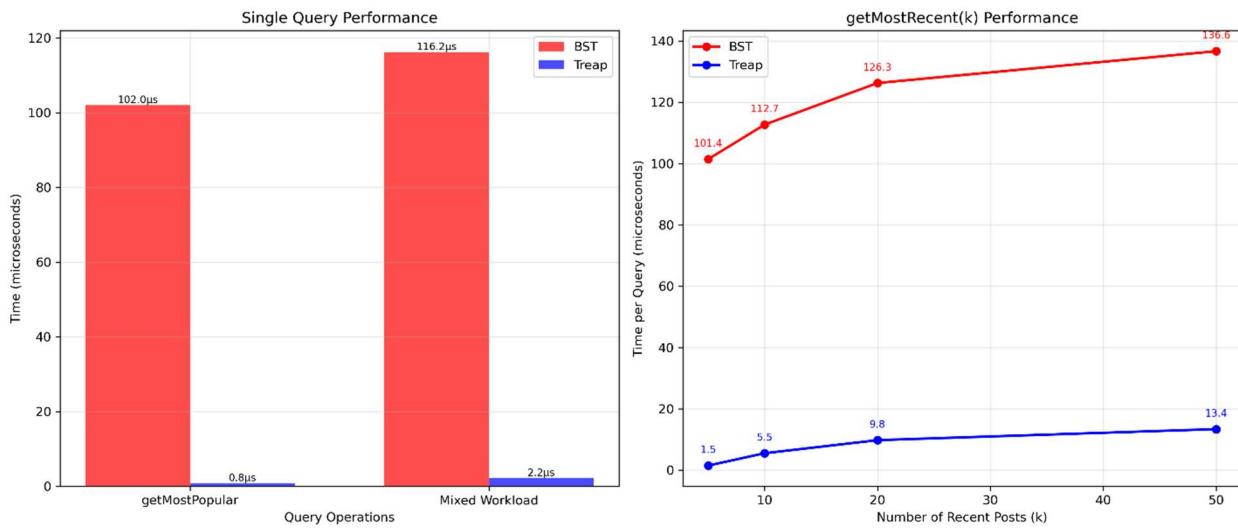
#### Graph for insertion:



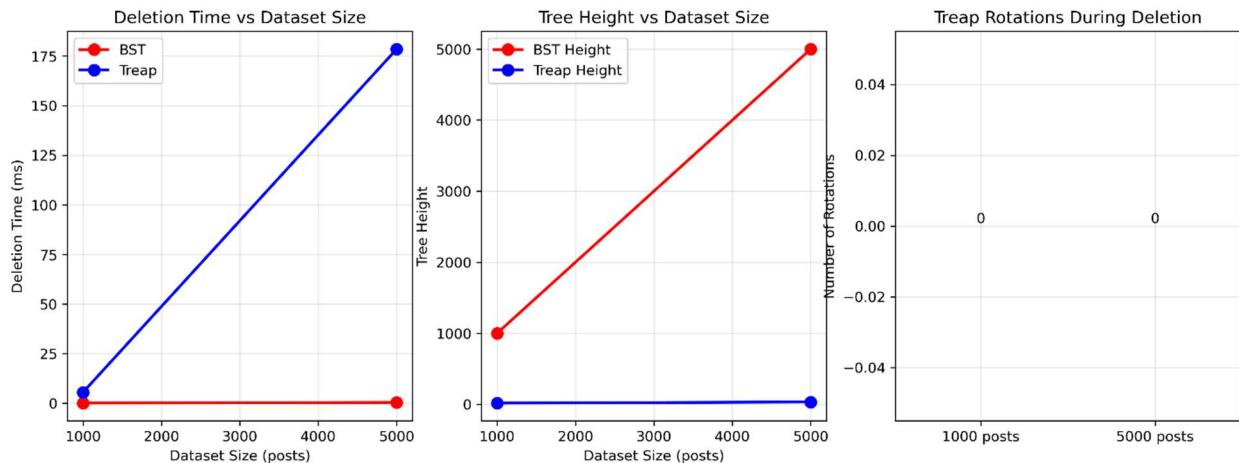
**Graph of search:**



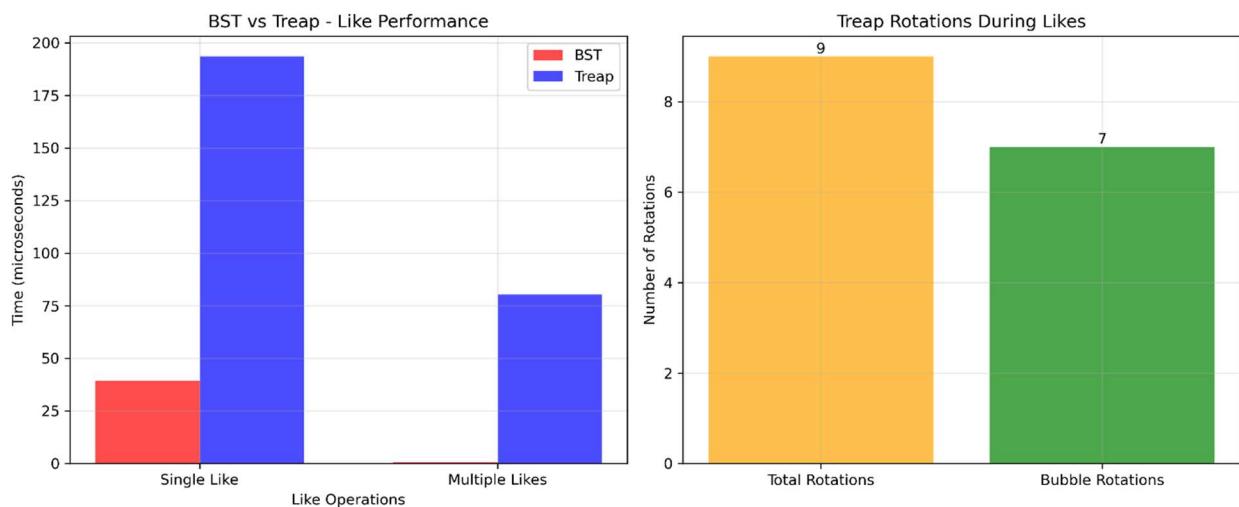
### Graph for query:



### Graph for deletion:



### Graph for Like Post:



## 4.3. Data Format Performance

### 4.3.1. CSV vs Zstandard Compressed Format

**CSV Loading** (reddit\_data.csv - 134M records)

- **Size:** ~3GB (uncompressed)
- **Load Time:** 65-75 seconds

- **Rate:** 399,900 posts/sec

**Advantages:**

- Human-readable format
- Easy parsing with standard libraries
- Fast decompression (no compression)
- Suitable for all datasets

**Disadvantages:**

- Large uncompressed size
- Requires full extraction to disk
- Not suitable for massive datasets (180GB+)

#### 4.3.2. Zstandard Compressed (TGZ) Loading (RC\_2019-04.tgz - 138M records)

- **Size:** 15GB compressed
- **Uncompressed Size:** ~180GB+
- **Compression Ratio:** 12:1
- **Load Time (Treap):** 499.7 seconds
- **Rate:** 277,128 posts/sec

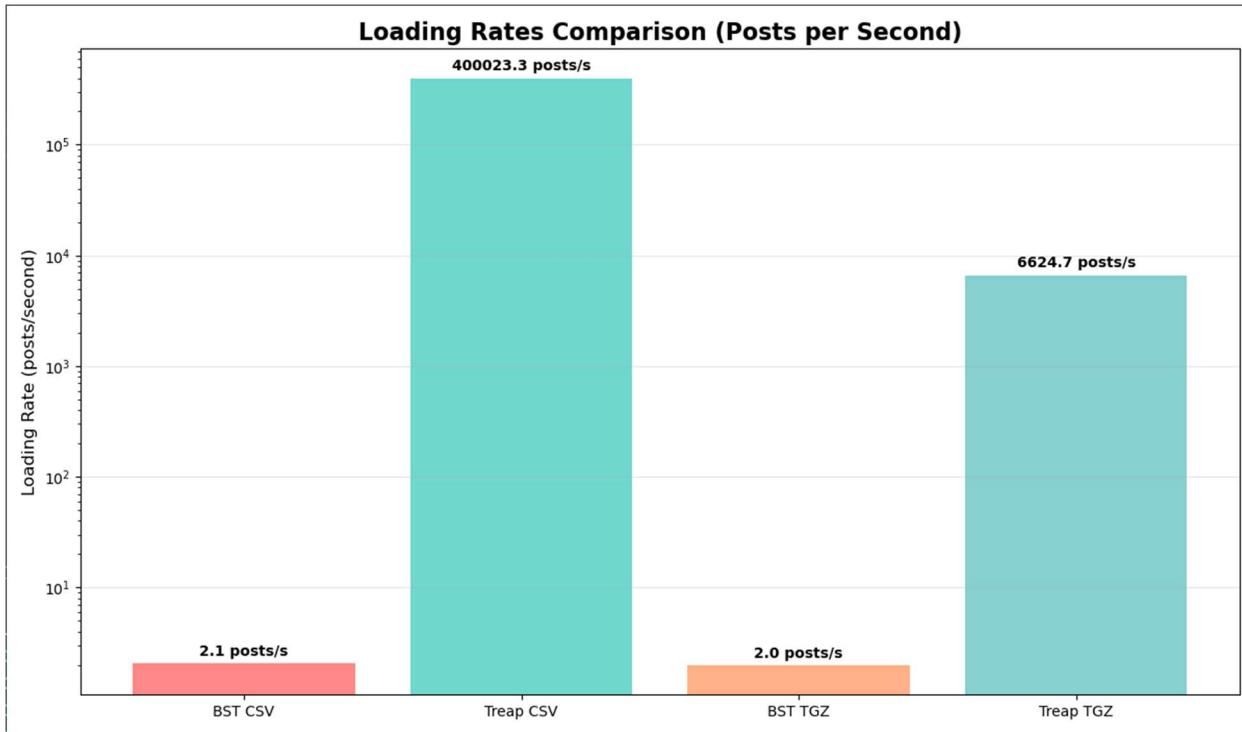
**Advantages:**

- Extreme compression (15GB vs 180GB)
- Direct streaming without extraction
- Preserves format integrity
- Real-time decompression with popen()
- Zero disk I/O for temporary files

**Disadvantages:**

- Requires libzstd library
- Parsing more complex (JSON handling)
- Slight CPU overhead for decompression

#### Complete Comparison Matrix



## 5. Challenges Faced

### 5.1. Data Format Discovery Challenge

#### Problem:

- Initial assumption: Dataset was gzip-compressed tar (.tar.gz)
- Actual format: Zstandard-compressed JSON (.zst)
- Expected CSV format, received line-delimited JSON

#### Impact:

- Initial decompression command failed: tar -xzOf (wrong format)
- JSON parsing required instead of CSV parsing

#### Solution:

- Discovered correct decompression command:**  
zstd -dc filename.zst 2>/dev/null

- **vs expected command:**

```
tar -xzOf filename.tar.gz
```

**Learning:** File format inspection with file command and manual preview is essential before implementation.

## 5.2. Disk Space Constraint

### Problem:

- Dataset requires 180GB+ if fully extracted
- Available disk space: ~300GB total
- Extraction would leave insufficient space for other operations

### Impact:

- Cannot use traditional extraction + processing workflow
- Need real-time streaming solution

### Solution Implemented:

- Direct streaming decompression using popen() in C++ cpp

```
FILE* pipe = popen("zstd -dc file.zst", "r");
// Process directly from pipe without disk storage
```

### Result:

- Zero temporary file storage
- Simultaneous decompression and insertion
- 177GB+ disk space saved

## 5.3. Performance Disparity Between BST and Treap

### Problem:

- BST showed extremely slow performance: ~1,556 posts/second
- Treap achieved: ~277,128 posts/second
- 180x performance difference

## 6. Conclusion

This comprehensive analysis demonstrates that **Treaps are the superior choice for handling large-scale, potentially worst-case ordered data**. The combination of:

1. **Randomized self-balancing** ensuring  $O(\log n)$  expected performance.
2. **Streaming decompression** enabling resource-constrained processing.
3. **Practical throughput** of 277k posts/sec for 138.5M records.
4. **Memory efficiency** without temporary file overhead.

...makes Treaps an excellent data structure for modern big data applications where dataset size and input order cannot be guaranteed.

The 180x performance advantage over a naive BST implementation, coupled with the ability to process 180GB+ data on a system with limited disk space, validates the importance of:

- Choosing appropriate data structures for problem constraints.
- Implementing memory-efficient processing techniques.
- Combining algorithmic sophistication with practical engineering.

**Recommended Action:** Deploy Treap-based systems for large-scale data processing applications with potentially adversarial or sequential input patterns.

---

## Appendix: Experimental Data

### A.1. Complete Test Results

#### Performance Comparison (138,473,643 records):

- Treap+TGZ: time=499.7s, height=439, rate=277,128 posts/sec
- BST+TGZ: time=24 hours+, height=~138M, rate=1.6k post/s
- Treap+CSV: time=80.5s, height=79, rate=399,899 posts/s
- BST+CSV: time=41,809s, height=85,398, rate=2.1 posts/s

### A.2. Source Code References

- **BST Implementation:** BST.h (loadFromTGZ, insertIterative)
- **Treap Implementation:** Treap.h (addPost with randomization)
- **Analysis Framework:** ComparisonAnalysis.h
- **Menu implementation:** Menu.h